

Assignment1__Seshadri

Sri Seshadri

4/7/2018

Section 2.8, Question 2.1 page 59 - Effect of transformation on time series data.

a) Monthly total of people on unemployment benefits in Australia (Jan 1956 - July 1992)

Figure 1 shows the time series of number of people in unemployment benefits in Australia by month. The time series is affected by

- * Population growth over time
- * External factors like:
 - * state of the economy
 - * the benefit provided by the government.

It would be useful to normalize the data by population, to get the percent unemployed of the total population. Then if need be a transformation on the normalized data can be made.

b) Monthly total accidental deaths in the United States

The top plot in figure 2 shows the total accidental deaths by month in the US. There is seasonality in the data, where the total accidents peaking at July. The variation in the seasonality may be mitigated by normalizing the totals by dividing by the number of days in the month.

The middle plot in figure 2 shows the normalized total by days in the months; i.e. Monthly average accidental deaths per day. We see that there is some smoothing of the raw data. While there is not much variation in seasonality, there is interest in further making the size of the seasonal variation equal across seasons.

The bottom chart of figure 2 shows a box-cox transformation of the average accidental deaths. Its seen that there isn't much affect as expected.

c) Quarterly production of bricks (in millions) at Portland, Australia

The top chart in Figure 3 shows the time series plot of quarterly brick production at Portland Australia. The time series exhibits an increasing trend and seasonality. The variation increases with time / levels. Box-Cox transformation is appropriate for this case. The ideal lambda for the data was 0.255. The bottom chart of figure 3 shows the transformed data. The variation is better across time, of course the huge downward spikes is not fully mitigated.

Question 2.2 Page 60

Time series modeling of Dow Jones index

In this section the Dow jones industrial average is modeled as a time series. The drift method is used to fit the time series and forecasted for the next 10 periods. Figure 4 shows the time series plot with forecast from drift,mean and naive mean methods. Table 1 shows the forecasts from drift method and the blue line in fig 4 plots the forecasts. The drift method forecast is nothing but the extension of the line that joins the first

Total people on unemployment benefits in Australia




Figure 1: Monthly number of people on unemployment benefits in Australia

and the last point. The slope of the line is the difference between the last and the first point divided by the number of data points. The grey dashed line in fig 4 shows the fitted drift for the Dow Jones data.

Table 1: Drift method forecast for 10 periods

Time	index
79	121.3636
80	121.4973
81	121.6309
82	121.7645
83	121.8982
84	122.0318
85	122.1655
86	122.2991
87	122.4327
88	122.5664

Comparison between models

The drift method is compared with other models like the Mean of the time series, the naive method and the seasonal naive method. The accuracy methods are compared in table 2. Since the time series is set up as a daily closing index. The seasonal forecast for the next point is identical as the previous point (same season as the last point). Hence the Naive and Seasonal Naive forecasts are the same for this scenario.

The Drift method is the best performing model amongst the competing naive models, based on MAE / MAPE. This is not surprising as we see an increasing trend in the Dow Jones index. The other methods use the mean or the last value of the time series as the forecast. The naive models do not account for the increasing trend. The Drift method does.

Table 2: Accuracy metrics of benchmark models

	ME	RMSE	MAE	MPE	MAPE	MASE	ACF1
Drift	0.000	0.424	0.325	-0.001	0.280	0.952	0.422
Mean	0.000	5.471	5.104	-0.222	4.400	14.938	0.985
Naive	0.134	0.445	0.342	0.114	0.294	1.000	0.422

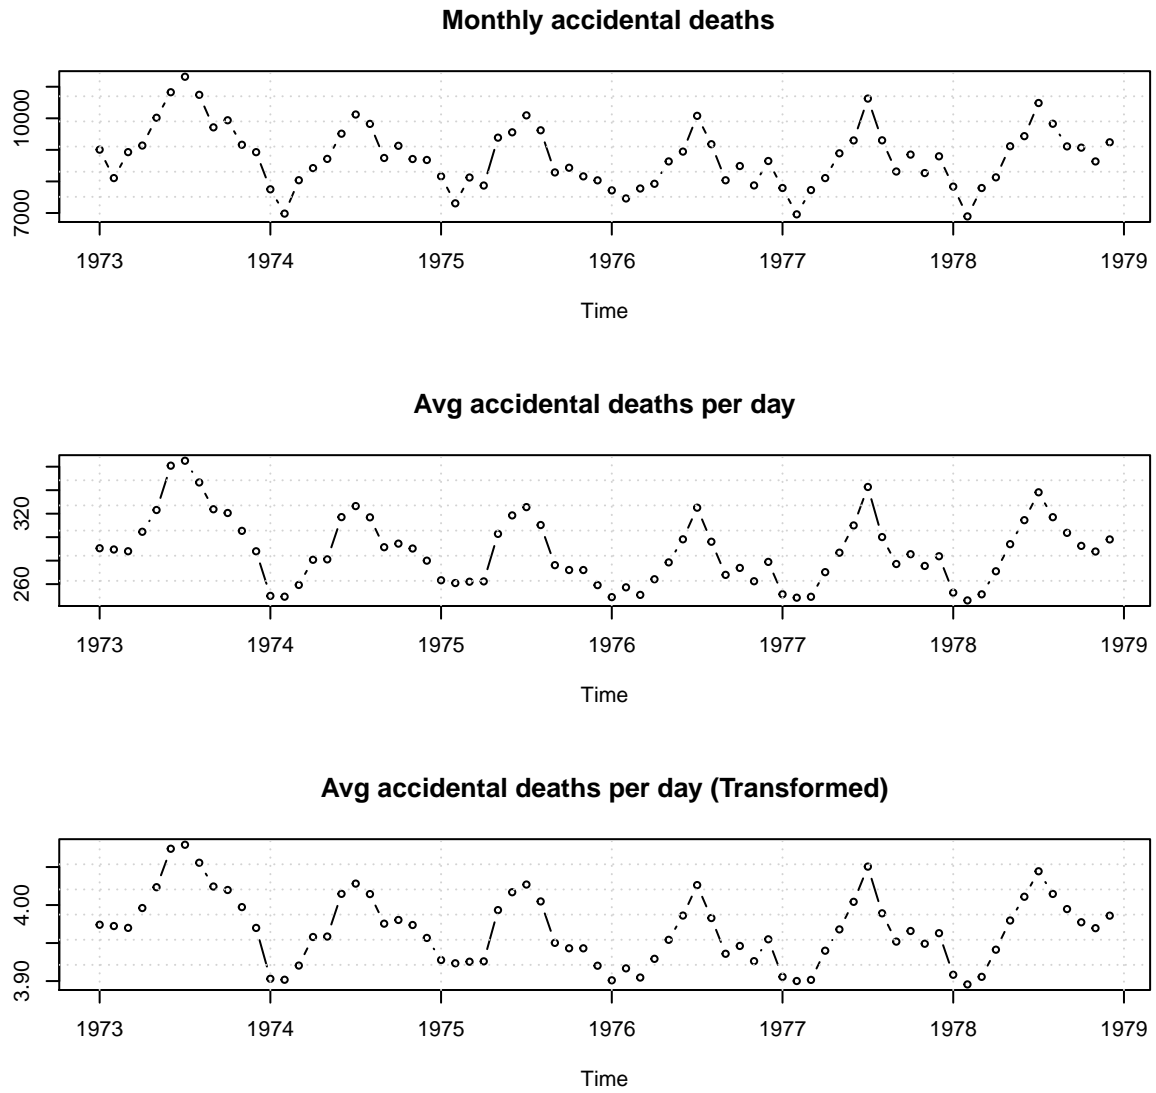


Figure 2: Top: Total accidental deaths by month in the US, Middle: Monthly Average Accidental deaths per day, Bottom: Transformed monthly average accidental deaths per day





Figure 3: Top: Quarterly production of bricks (in millions) at Portland, Australia, Bottom : BoxCox transformed Quarterly production of bricks




Figure 4: Time series modeling of Dow Jones industrial average

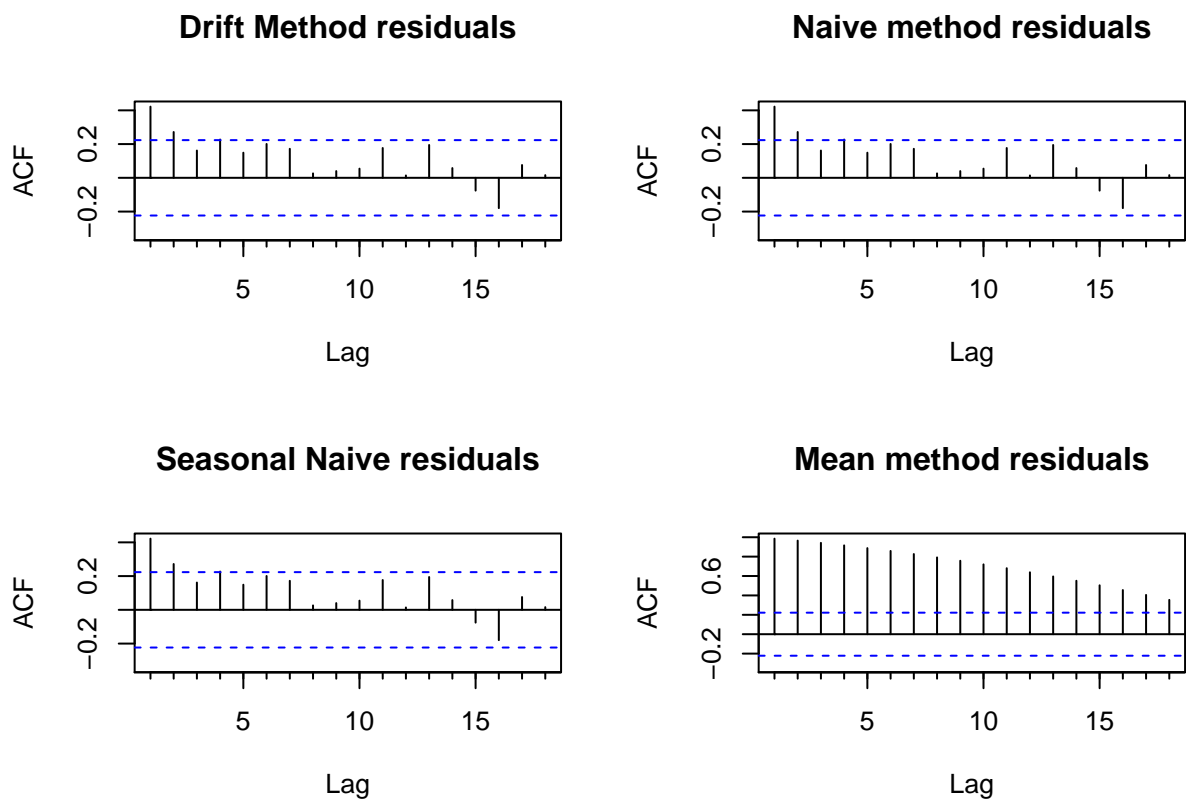
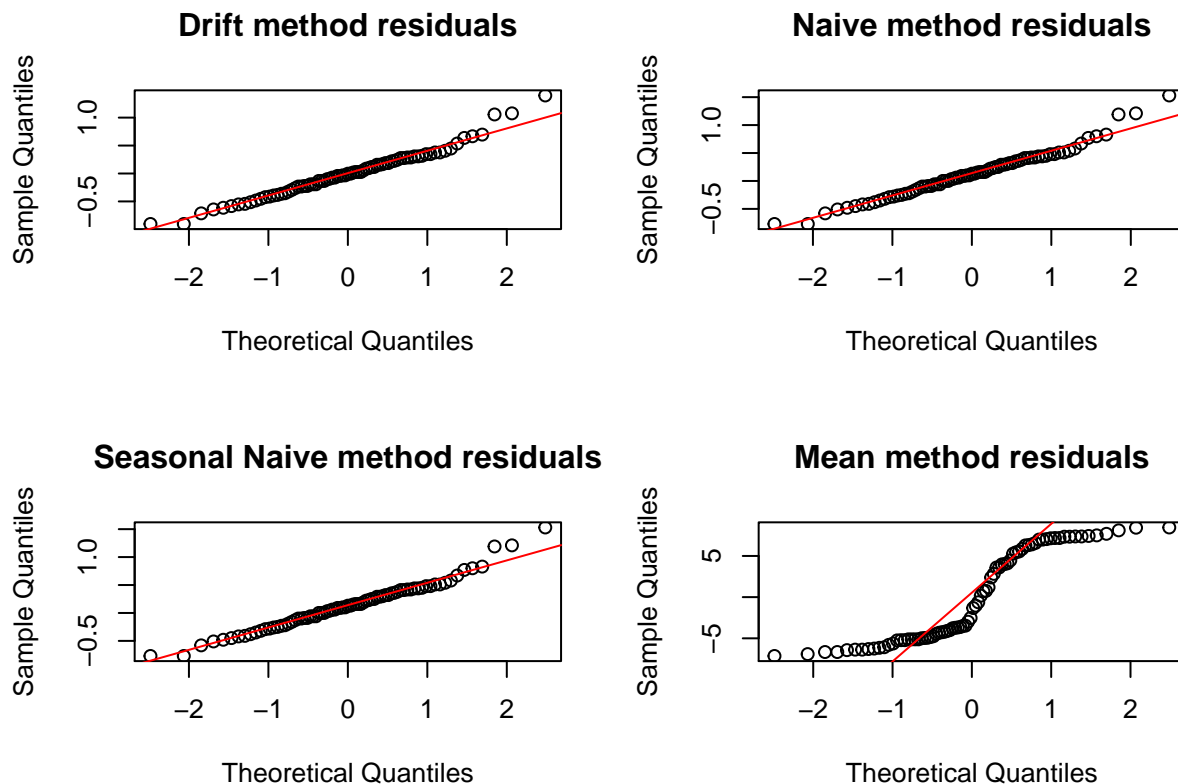


Figure 5: Autocorrelation plot of residuals



Question 2.3 IBM Stock prices

Figure 8 shows the time series plot of IBM stock price at close. The data is split into training and test set. The first 300 data points are retained as training set and the rest of the 69 data points are held out as test set.

Time series modeling of the training set

The Drift, naive and mean method modeling was trained on the training set and the accuracy statistics are shown in table 3. Since the distribution of the data is bimodal, BoxCox transformation was deemed not appropriate. Figure 9 shows the residual analysis. Both the naive and the drift methods are very close in their forecasts.

Table 3: Forecast Accuracy

	ME	RMSE	MAE	MPE	MAPE	MASE	ACF1
Drift	0.00	7.30	5.13	-0.03	1.12	1.01	0.14
Naive	-0.28	7.30	5.10	-0.08	1.12	1.00	0.14
Mean	0.00	73.62	58.72	-2.64	13.03	11.52	0.99

Figure 10 shows the actuals and the forecasts along with the time series of the forecast errors. Both the naive and drift methods are very identical.

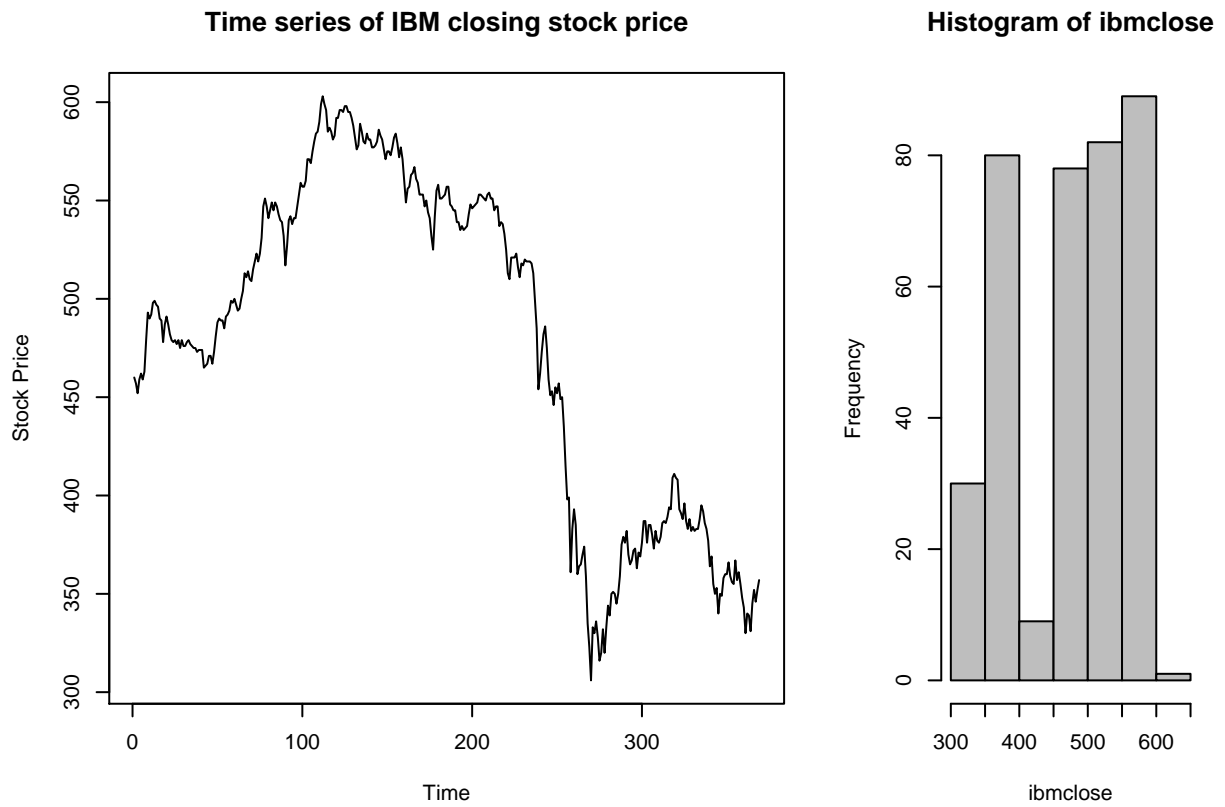
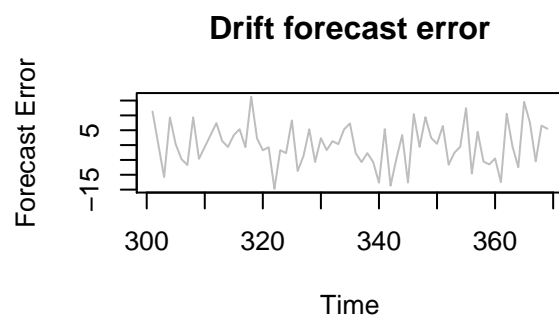
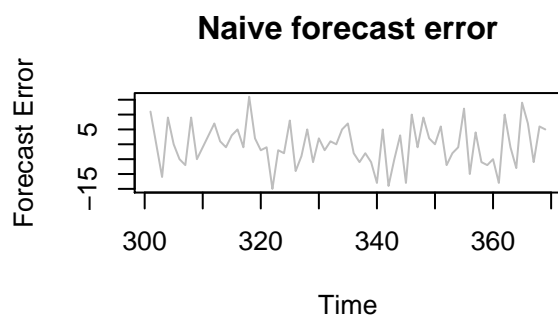




Figure 6: IBM Stock price at close



Drift Residuals



Naive Residuals




Figure 7: Residual Analysis

